

The image features the word "NUMPY" in a bold, dark blue, sans-serif font, centered within a white circular area. This circle is set against a background composed of three distinct color regions: a light blue area on the left, a light pink area on the right, and a dark blue area at the bottom. The white circle is positioned such that it overlaps the top and sides of the dark blue area, creating a central focus for the text.

**NUMPY**

# WHY NUMPY?

```
import numpy as np
```

## Python is slow for numerical computing

- Python itself is an interpreted language — it processes one instruction at a time.
- Doing heavy number crunching (e.g., multiplying millions of numbers) in plain Python loops is **very slow**.

That's where **NumPy** comes in.

## NumPy is mostly written in C

- The core of NumPy (ndarray objects and mathematical operations) is implemented in C (and some Fortran).
- C is a **compiled language**, meaning the code runs directly on your machine's processor → **much faster than Python loops**.

# CREATING ARRAYS

```
a = np.array([1, 2, 3])           # 1D
M = np.array([[1,2,3],[4,5,6]])  # 2D
zeros = np.zeros((3,3))
ones  = np.ones(4)
r = np.arange(0,10,2)            # like range()
lin = np.linspace(0,1,5)         # evenly spaced
I = np.eye(3)                   # identity
```

Exercise:

Create a 4x2 array of ones and  
a vector from 10 down to 1.

# ARRAY ATTRIBUTES

```
a.shape    # tuple of dims
a.ndim     # number of dimensions
a.size     # total elements
a.dtype    # data type
```

# INDEXING AND SLICING

## Basic Indexing

```
v = np.array([10,20,30,40])  
v[0]      # first  
v[-1]     # last  
v[1:3]    # slice (view)
```

## 2D Indexing

```
M[1,2]    # row 1 col 2  
M[:,1]    # all rows, col 1  
M[0]      # first row
```

## Copying an array

```
s = v[1:3]  
s[0] = 999    # changes original v  
copy = v[1:3].copy() # make a copy instead
```

Exercise: slice a 3×4 matrix to get the middle two columns and modify one value.

# VECTORIZED OPERATIONS

```
a = np.array([1,2,3])
b = np.array([4,5,6])
a + b          # [5,7,9]
a * 2          # [2,4,6]
a * b          # elementwise [4,10,18]
a @ b          # dot product (or np.dot(a,b))
```

- Much faster than Python loops.  
Exercise: compute column-wise mean of a 2D array using vectorized ops.

# BROADCASTING

Compatibility rules:

- dimensions either equal or one is 1;
- NumPy aligns trailing dims.

```
A = np.ones((2,3))  
v = np.array([1,2,3])    # shape (3,)   
A + v    # v broadcasts across rows => each row + v
```

```
a = np.array([[1], [2], [3]])    # shape (3,1)  
b = np.array([10, 20, 30])      # shape (3,)
```

```
print(a + b)
```

```
import numpy as np
```

```
arr = np.array([1, 2, 3, 4])  
print(arr + 10)    # [11 12 13 14]
```

Exercise: add a (3,1) column vector to a (3,4) matrix

# AGGREGATION AND AXIS

```
data = np.array([[1, 2, 3],
                 [4, 5, 6]])

print(data.sum())           # 21  (all elements)
print(data.sum(axis=0))     # [5 7 9]  (column-wise)
print(data.sum(axis=1))     # [6 15]  (row-wise)


print(data.mean())         # 3.5
print(data.min(axis=1))    # [1 4]
print(data.max(axis=0))    # [4 5 6]
```

Exercise: compute average per row and per column for a sample matrix.

# LINEAR ALGEBRA

```
A = np.array([[1,2],[3,4]])  
np.linalg.inv(A)  
np.linalg.det(A)  
np.linalg.eig(A)
```

Exercise: compute eigenvalues of a simple matrix.





# PERFORMANCE COMPARISON

## PYTHON LIST VS. NUMPY ARRAY

9

Python talks to C via a “wrapper”

- NumPy provides a Python API (functions like `np.array`, `np.dot`, etc.).
- When you call a NumPy function in Python, under the hood it calls optimized C functions.
- This is possible because:
  - Python allows C extensions (modules written in C can be imported in Python).
  - NumPy is such a C extension.

```
import time
n = 1_000_000
L = list(range(n))
start = time.time()
[s*2 for s in L]
print("list time", time.time()-start)

import numpy as np
a = np.arange(n)
start = time.time()
a*2
print("numpy time", time.time()-start)
```

# PERFORMANCE COMPARISON PYTHON LIST VS. NUMPY ARRAY

## Vectorization vs SIMD

- NumPy's core vectorization = **looping in C over raw memory buffers** (fast).
- Some NumPy builds (e.g., with Intel MKL, OpenBLAS, or SIMD instructions like AVX) **use CPU vector instructions**.
  - This means the CPU itself processes multiple elements in one instruction.
  - Example: adding 4 doubles in one CPU cycle instead of one at a time.
- That's why performance can differ depending on your NumPy build.

10

