

Sorting Algorithms

The logo for AIT (Asian Institute of Technology) is displayed within a white circular frame. It consists of the letters 'A', 'I', and 'T' in a stylized, blue, sans-serif font. The 'A' is formed by two parallel diagonal lines meeting at the top, and the 'I' and 'T' are solid vertical strokes with horizontal bars at the top.

AIT



What is Sorting?

Sorting is the process of reordering a collection of items into a specific order, typically ascending or descending.

While we often talk about sorting simple numbers, in practice, we usually sort records based on a specific **key**. For example, sorting a list of students by their ID number or by last name.

- **Input:** An unsorted list $(a_0, a_1, \dots, a_{n-1})$
- **Output:** A sorted permutation $(a'_0, a'_1, \dots, a'_{n-1})$ such that $a'_0 \leq a'_1 \leq \dots \leq a'_{n-1}$



A Key Constraint: In-Place Sorting

Memory usage is a critical factor in algorithm design.

- An **in-place** sorting algorithm uses a constant amount of extra memory, denoted as $\Theta(1)$. It sorts the elements within the original array, perhaps using a few extra variables for swaps.
- Other algorithms require allocating a significant amount of additional memory, often a second array of the same size ($\Theta(n)$ *memory*).

For memory-constrained environments, **in-place algorithms are strongly preferred.**



Common Sorting Strategies

Sorting algorithms can be categorized by their core strategy.

Understanding these helps in grasping how different algorithms work.

- **Insertion:** Build the final sorted array one item at a time.
- **Exchanging:** Systematically swap elements that are in the wrong order.
- **Selection:** Repeatedly find the next smallest element and move it to its final position.
- **Merging:** Divide the array, sort the sub-arrays, and then merge them back together.
- **Distribution:** Distribute elements into buckets based on their values.



The Big Picture: Runtime Classes

The efficiency of sorting algorithms generally falls into one of three categories. This is the most important factor in choosing an algorithm.

- $\Theta(n^2)$: Simple to implement but slow on large datasets. (e.g., Insertion Sort, Bubble Sort)
- $\Theta(n \log n)$: The standard for efficient, general-purpose sorting. (e.g., Heap Sort, Merge Sort, Quicksort)
- $\Theta(n)$: Extremely fast but can only be used under special assumptions about the data (e.g., the items are integers within a known, small range). (e.g., Bucket Sort, Radix Sort)



The Theoretical Speed Limit

For any sorting algorithm that relies on comparing elements to each other, it's impossible to do better than $\Omega(n \log n)$ time in the average and worst cases.

Why?

There are $n!$ possible permutations of n items. Each comparison can only cut the number of remaining possibilities in half. The number of comparisons needed to distinguish between all $n!$ outcomes is at least $\log_2(n!)$, which can be shown to be $\Omega(n \log n)$.

This means that algorithms like Heap Sort and Merge Sort are, in a theoretical sense, **asymptotically optimal**.



So, Which Algorithm is "Best"?

There is **no single best sorting algorithm** for all situations.

The optimal choice depends on the specific context:

- How large is the dataset?
- Is the data already "nearly sorted"?
- Are there strict memory limitations (requiring an in-place sort)?
- Is the algorithm's worst-case performance a critical concern?



The "Anti-Best" Algorithm (For a Laugh 😂)

To appreciate good algorithms, it helps to see a comically bad one.

Bogosort (aka "Permutation Sort"):

1. Randomly shuffle the list.
 2. Check if the list is sorted.
 3. If not, go back to step 1.
- **Best-Case Runtime:** $\Theta(n)$ (You get lucky on the first shuffle).
 - **Average-Case Runtime:** $\Theta(n * n!)$. This is astronomically slow.
 - **Worst-Case Runtime:** Unbounded. It might never finish!



Measuring "Unsortedness"

How can we quantify how "mixed up" a list is? A list can be slightly out of order or completely random.

We need a formal way to measure this, which can help predict the performance of certain algorithms.



Defining an Inversion

An inversion is a pair of elements in a list that are out of their natural sorted order.

Formally, for a list a , a pair of indices (j, k) forms an inversion if:

$$j < k \text{ but } a[j] > a[k]$$

- A perfectly sorted list has 0 inversions.
- A reverse-sorted list has the maximum possible number of inversions.



Inversions by Example

Let's find the inversions in the list $(1, 3, 5, 4, 2, 6)$.

We look for pairs where a larger number appears before a smaller one:

- $(3, 2)$
- $(5, 4)$
- $(5, 2)$
- $(4, 2)$

This list has **4 inversions**.



Why Do Inversions Matter?

The number of inversions is directly related to the runtime of sorting algorithms that work by swapping adjacent, out-of-order elements (like **Bubble Sort** and **Insertion Sort**).

Each adjacent swap can remove at most one inversion. Therefore, the runtime of such algorithms is at least proportional to the number of inversions in the initial list. For a "nearly sorted" list with few inversions, these algorithms can be very fast.



How Many Inversions to Expect?

For a list of size n , there are $\binom{n}{2} = (n(n-1))/2$ total pairs of elements.

If the list is randomly shuffled, we expect about half of these pairs to be inversions.

Expected Inversions in a Random List: $\approx 1/2 * (n(n-1))/2 = (n^2 - n)/4$, which is $\Theta(n^2)$.

This tells us that an average, random list is highly unsorted, which is why $\Theta(n^2)$ algorithms are a natural starting point.



Analyzing Lists with Inversions

Consider three lists of 20 numbers, where a random list is expected to have about 95 inversions.

- **List 1 (Nearly Sorted):** *1 16 12 26 25 ...*
 - Has only **13 inversions**. An algorithm like Insertion Sort would be very fast here.
- **List 2 (Mostly Sorted with Outliers):** *1 17 21 42 24 ... 57 23 ...*
 - Also has **13 inversions**, but the out-of-place items are far from their correct positions.
- **List 3 (Randomly Sorted):** *22 20 81 38 95 ...*
 - Has **100 inversions**, very close to the expected random value.



Summary

- **Sorting** is the process of reordering a list based on a key.
- **In-place** algorithms ($\Theta(1)$ *extra memory*) are highly desirable.
- The main performance classes are $\Theta(n^2)$ (*simple*) and $\Theta(n \log n)$ (*efficient*).
- A fundamental **lower bound** for comparison-based sorting is $\Omega(n \log n)$.
- An **inversion** is a pair of out-of-order elements and serves as a formal measure of how unsorted a list is. The number of inversions can directly impact the performance of certain sorting algorithms.