

# Memory, Data, & Addressing I

**Rabbit**



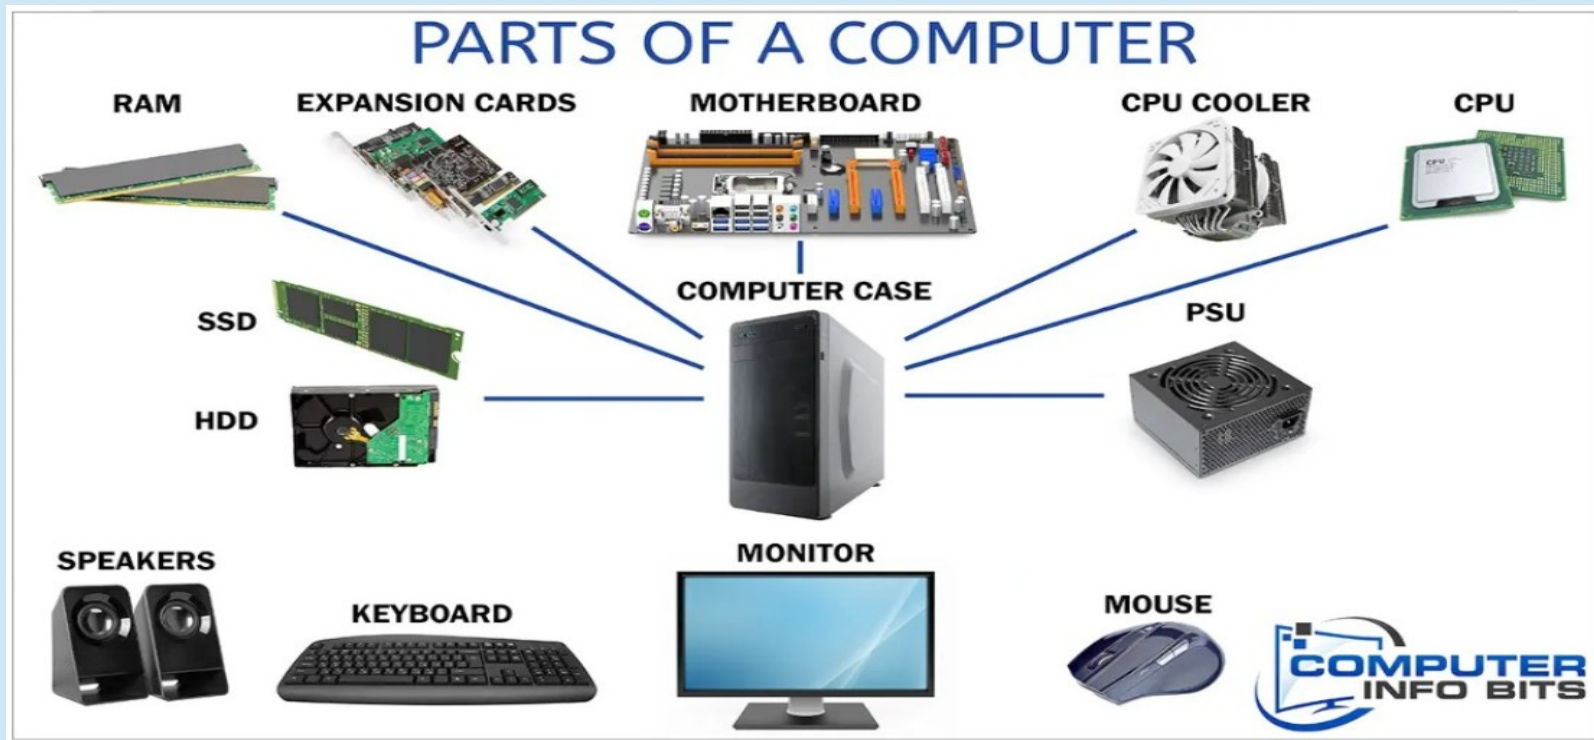
**Rabbyte**




**Rabword**



## Computer Parts




## Hardware: Physical View




## Hardware: Logical View

*Hooray,  
abstraction!*




Hardware:




- The **CPU** executes instructions
- **Memory** is where data (including instructions) is stored
- How is data encoded? ***Binary encoding!***

Hardware:



- To execute an instruction, the CPU must:
  1. Fetch instruction
  2. (if applicable) Fetch data needed for the instruction
  3. Perform the computation
  4. (if applicable) Write the result back to memory

## Hardware: Today's focus



- We'll start by focusing on memory
  - a. How does the CPU know where to find its data?
  - b. How are common data types encoded?
  - c. How can we use C to manipulate data?

## Review Questions (pt 1)

1. By looking at the bits stored in memory, I can tell what a particular 4 bytes is used to represent.

**A) True**                      **B) False**

2. We can fetch a piece of data from memory as long as we have its address.

**A) True**                      **B) False**



## Review Questions (pt 2)

3. Which of the following bytes have a most-significant bit (MSB) of 1?

- A)** 0x63      **B)** 0x90      **C)** 0xCA      **D)** 0xD

4. Consider the following declared variables:

```
int x = 351;
```

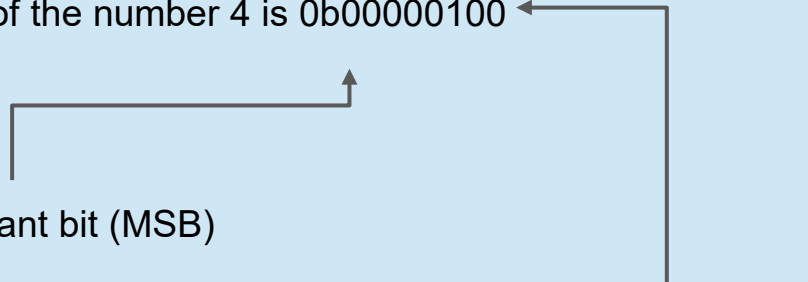
How much space in memory does variable p take up (on a 64-bit machine)?

- A)** 1 byte    **B)** 2 bytes    **C)** 4 bytes    **D)** 8 bytes

## Fixed-length Binary

- Because storage is finite, everything is stored as “fixed” length
  - Data is moved and manipulated in fixed-length chunks
  - Multiple fixed lengths (e.g., 1 byte, 4 bytes, 8 bytes)
  - Leading zeros now must be included up to “fill out” the fixed length

Example: the “eight-bit” (1-byte) representation of the number 4 is 0b00000100



Most significant bit (MSB)


Least significant bit (LSB)

## Bits and Bytes and Things

- **Useful Fact:**  $n$  bits can represent up to  $2^n$  things
  - If we want to represent  $x$  things, we need  $n$  bits, where  $2^n \geq x$   
Example: how many bits would we need to represent the letters a-z?  
26 letters, so we need 5 bits ( $2^5 = 32 > 26$ )
- Sometimes (oftentimes?) the “things” we are representing are **bytes**!

## Addresses in Memory

- You can think of memory as a single, large array of bytes, each with a unique **address** (index)
  - Addresses are fixed-width
  - Amount of addressable memory = **address space**
- If addresses are  $a$  bits long, how many addresses are there?
  - So how big is the address space?



## Machine “Words”

- word size = address size
- Most modern systems use **64-bit (8-byte) words**
  - Address space = 2 addresses, each represents a byte of data
  - $2^{64}$  bytes =  $1.8 \times 10^{19}$  bytes = 18 EB (exabytes)
    - (Side note: your computer does not actually have this much memory! We'll talk about this more later)


## Address of Multibyte Data

- Addresses still specify locations of bytes, but we can choose to view memory as a series of fixed-size chunks instead
  - Addresses of successive chunks differ by data size
- The address of any chunk of memory is the lowest address of the chunk
  - To specify a chunk, need *both* its address and size
  - Typically **aligned**, meaning their starting addresses are multiples of the data size

64-bit data	32-bit data	Bytes	Addr. (hex)
Addr = 0000	Addr = 0000		0x00
			0x01
			0x02
			0x03
	Addr = 0004		0x04
			0x05
			0x06
			0x07
Addr = 0008	Addr = 0008		0x08
			0x09
			0x0A
			0x0B
	Addr = 0012		0x0C
			0x0D
			0x0E
			0x0F

## Byte Ordering


- How should bytes within a piece of data be ordered in memory?
  - Similar to writing in human languages
    - Ex: English reads left-right, but Arabic reads right-left
- By convention, ordering of bytes is called *endianness*
  - Two options: **big-endian** and **little-endian**
  - Reference to *Gulliver's Travels*: tribes cut their eggs on different sides



## Endianness

- **Big-endian** (SPARC, z/Architecture)
  - Least-significant byte at the highest address
- **Little-endian** (x86, x86-64)
  - Least-significant byte at the lowest address
- **Bi-endian** (ARM, PowerPC)
  - Endianness can be specified as either big or little

Example: 4-byte data 0xA1B2C3D4 at address 0x100





## Polling Question

- We store the value 0x 01 02 03 04 as a *word* at address 0x100 in a big-endian, 64-bit machine
- What is the **byte of data** stored at address 0x104?

A) 0x04

B) 0x40

C) 0x01

D) 0x10

E) We're lost...

0x100	0x101	0x102	0x103	0x104	0x105	0x106	0x107

## Endianness (pt 2)


- Endianness only applies to data storage
- Often, a programmer can ignore endianness because it is handled for them
  - Bytes wired into correct place when reading and storing from memory (hardware)
  - Compiler and assembler generate correct behavior (software)
- Endianness still shows up:
  - Logical issues: accessing different amount of data than how you stored it (e.g., store int, access byte as a char)
  - Need to know exact values to debug memory errors
  - Manual translation of machine code

## Data Representations in C

Java Data Type	C Data Type	32-bit (old)	64-bit
boolean	bool	1	1
byte	char	1	1
char		2	2
short	short/short int	2	2
int	int	4	4
float	float	4	4
	long/long int	4	8
double	double	8	8
long	long long/long long int	8	16
	long double	8	16
<b>(reference)</b>	<b>pointer (*)</b>	<b>4</b>	<b>8</b>

## A Picture of Memory (64-bit view)

- A 64-bit (8-byte) aligned view of memory, big endian
  - Each cell is one byte
  - Each row is composed of 8 bytes
  - The labels on the left are the starting addresses of each row



## Summary

- Memory is a long, *byte-addressed* array
  - **Word size** bounds the size of the address and memory (**address space**)
  - Different data types use different number of bytes
  - Address of multi-byte piece of memory given by the lowest address
- **Endianness** determines memory storage order for multi-byte data
  - Least significant byte in lowest (**little-endian**) or highest (**big-endian**) address of memory chunk