

GCC Optimizations

Optimization

Most of what you need to do with optimization can be summarized by:

- 1) If doing something seldom and only on small inputs, do whatever is simplest to code, understand, and debug
- 2) If doing things a lot, or on big inputs, make the primary algorithm's Big-O cost reasonable
- 3) Let gcc do its magic from there**
- 4) Optimize explicitly as a last resort

GCC Optimization

- Today, we'll be comparing two levels of optimization in the gcc compiler:
 - `gcc -O0` // mostly just literal translation of C
 - `gcc -O2` // enable nearly all reasonable optimizations
 - (we also use `-Og`, like `-O0` but more debugging friendly)
- There are other custom and more aggressive levels of optimization, e.g.:
 - `-O3` //more aggressive than `O2`, trade size for speed
 - `-Os` //optimize for size
 - `-Ofast` //disregard standards compliance (!!)
- Exhaustive list of gcc optimization-related flags:
 - <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>

GCC Optimizations

- Constant Folding
- Common Sub-expression Elimination
- Dead Code
- Strength Reduction
- Code Motion
- Tail Recursion
- Loop Unrolling

Constant Folding

Constant Folding pre-calculates constants at compile-time where possible.

```
int seconds = 60 * 60 * 24 * n_days;
```

Constant Folding

```
int fold(int param) {  
    char arr[5];  
    int a = 0x107;  
    int b = a * sizeof(arr);  
    int c = sqrt(2.0);  
    return a * param + (a + 0x15 / c + strlen("Hello") * b - 0x37) / 4;  
}
```

Constant Folding: Before (-00)

00000000000011b9 <fold>:

11b9:	55	push	%rbp	
11ba:	48 89 e5	mov	%rsp,%rbp	
11bd:	41 54	push	%r12	
11bf:	53	push	%rbx	
11c0:	48 83 ec 30	sub	\$0x30,%rsp	
11c4:	89 7d cc	mov	%edi,-0x34(%rbp)	
11c7:	c7 45 ec 07 01 00 00	movl	\$0x107,-0x14(%rbp)	
11ce:	8b 45 ec	mov	-0x14(%rbp),%eax	
11d1:	48 98	cltq		
11d3:	89 c2	mov	%eax,%edx	
11d5:	89 d0	mov	%edx,%eax	
11d7:	c1 e0 02	shl	\$0x2,%eax	
11da:	01 d0	add	%edx,%eax	
11dc:	89 45 e8	mov	%eax,-0x18(%rbp)	
11df:	48 8b 05 2a 0e 00 00	mov	0xe2a(%rip),%rax	# 2010 <_IO_stdin_used+0x10>
11e6:	66 48 0f 6e c0	movq	%rax,%xmm0	
11eb:	e8 b0 fe ff ff	callq	10a0 <sqrt@plt>	
11f0:	f2 0f 2c c0	cvttsd2si	%xmm0,%eax	
11f4:	89 45 e4	mov	%eax,-0x1c(%rbp)	
11f7:	8b 45 ec	mov	-0x14(%rbp),%eax	
11fa:	0f af 45 cc	imul	-0x34(%rbp),%eax	
11fe:	41 89 c4	mov	%eax,%r12d	
1201:	b8 15 00 00 00	mov	\$0x15,%eax	
1206:	99	cld		
1207:	f7 7d e4	idivl	-0x1c(%rbp)	
120a:	89 c2	mov	%eax,%edx	
120c:	8b 45 ec	mov	-0x14(%rbp),%eax	
120f:	01 d0	add	%edx,%eax	
1211:	48 63 d8	movslq	%eax,%rbx	
1214:	48 8d 3d ed 0d 00 00	lea	0xded(%rip),%rdi	# 2008 <_IO_stdin_used+0x8>
121b:	e8 20 fe ff ff	callq	1040 <strlen@plt>	
1220:	8b 55 e8	mov	-0x18(%rbp),%edx	
1223:	48 63 d2	movslq	%edx,%rdx	
1226:	48 0f af c2	imul	%rdx,%rax	
122a:	48 01 d8	add	%rbx,%rax	
122d:	48 83 e8 37	sub	\$0x37,%rax	
1231:	48 c1 e8 02	shr	\$0x2,%rax	
1235:	44 01 e0	add	%r12d,%eax	
1238:	48 83 c4 30	add	\$0x30,%rsp	
123c:	5b	pop	%rbx	
123d:	41 5c	pop	%r12	
123f:	5d	pop	%rbp	
1240:	c3	retq		

Constant Folding: After (-02)

```
00000000000011b0 <fold>:
 11b0: 69 c7 07 01 00 00      imul    $0x107,%edi,%eax
 11b6: 05 a5 06 00 00      add     $0x6a5,%eax
 11bb: c3                  retq
```

What is the consequence of this for you as a programmer? What should you do differently or the same knowing that compilers can do this for you?

GCC Optimizations

- Constant Folding
- **Common Sub-expression Elimination**
- Dead Code
- Strength Reduction
- Code Motion
- Tail Recursion
- Loop Unrolling

Common Sub-Expression Elimination

Common Sub-Expression Elimination prevents the recalculation of the same thing many times by doing it once and saving the result.

```
int a = (param2 + 0x107);  
int b = param1 * (param2 + 0x107) + a;  
return a * (param2 + 0x107) + b * (param2 + 0x107);
```

Common Sub-Expression Elimination

Common Sub-Expression Elimination prevents the recalculation of the same thing many times by doing it once and saving the result.

```
int a = (param2 + 0x107);  
int b = param1 * (param2 + 0x107) + a;  
return a * (param2 + 0x107) + b * (param2 + 0x107);  
// = 2 * a * a + param1 * a * a
```

```
000000000000011b0 <subexp>: // param1 in %edi, param2 in %esi  
    11b0: lea    0x107(%rsi),%eax    // %eax stores a  
    11b6: imul   %eax,%edi             // param1 * a  
    11b9: lea    (%rdi,%rax,2),%esi      // 2 * a + param1 * a  
    11bc: imul   %esi,%eax              // a * (2 * a + param1 * a)  
    11bf: retq
```

Common Sub-Expression Elimination

Why should we bother saving repeated calculations in variables if the compiler has common subexpression elimination?

- The compiler may not always be able to optimize every instance. Plus, it can help reduce redundancy!
- Makes code more readable!

GCC Optimizations

- Constant Folding
- Common Sub-expression Elimination
- **Dead Code**
- Strength Reduction
- Code Motion
- Tail Recursion
- Loop Unrolling

Dead Code

Dead code elimination removes code that doesn't serve a purpose:

```
if (param1 < param2 && param1 > param2) {  
    printf("This test can never be true!\n");  
}
```

```
// Empty for loop  
for (int i = 0; i < 1000; i++);
```

```
// If/else that does the same operation in both cases  
if (param1 == param2) {  
    param1++;  
} else {  
    param1++;  
}
```

```
// If/else that more trickily does the same operation in both cases  
if (param1 == 0) {  
    return 0;  
} else {  
    return param1;  
}
```

Dead Code: Before (-00)

00000000000011a9 <dead_code>:

11a9: 55	push	%rbp	
11aa: 48 89 e5	mov	%rsp,%rbp	
11ad: 48 83 ec 20	sub	\$0x20,%rsp	
11b1: 89 7d ec	mov	%edi,-0x14(%rbp)	
11b4: 89 75 e8	mov	%esi,-0x18(%rbp)	
11b7: 8b 45 ec	mov	-0x14(%rbp),%eax	
11ba: 3b 45 e8	cmp	-0x18(%rbp),%eax	
11bd: 7d 19	jge	11d8 <dead_code+0x2f>	
11bf: 8b 45 ec	mov	-0x14(%rbp),%eax	
11c2: 3b 45 e8	cmp	-0x18(%rbp),%eax	
11c5: 7e 11	jle	11d8 <dead_code+0x2f>	
11c7: 48 8d 3d 36 0e 00 00	lea	0xe36(%rip),%rdi	# 2004 <_IO_stdin_used+0x4>
11ce: b8 00 00 00 00	mov	\$0x0,%eax	
11d3: e8 68 fe ff ff	callq	1040 <printf@plt>	
11d8: c7 45 fc 00 00 00 00	movl	\$0x0,-0x4(%rbp)	
11df: eb 04	jmp	11e5 <dead_code+0x3c>	
11e1: 83 45 fc 01	addl	\$0x1,-0x4(%rbp)	
11e5: 81 7d fc e7 03 00 00	cmpl	\$0x3e7,-0x4(%rbp)	
11ec: 7e f3	jle	11e1 <dead_code+0x38>	
11ee: 8b 45 ec	mov	-0x14(%rbp),%eax	
11f1: 3b 45 e8	cmp	-0x18(%rbp),%eax	
11f4: 75 06	jne	11fc <dead_code+0x53>	
11f6: 83 45 ec 01	addl	\$0x1,-0x14(%rbp)	
11fa: eb 04	jmp	1200 <dead_code+0x57>	
11fc: 83 45 ec 01	addl	\$0x1,-0x14(%rbp)	
1200: 83 7d ec 00	cmpl	\$0x0,-0x14(%rbp)	
1204: 75 07	jne	120d <dead_code+0x64>	
1206: b8 00 00 00 00	mov	\$0x0,%eax	
120b: eb 03	jmp	1210 <dead_code+0x67>	
120d: 8b 45 ec	mov	-0x14(%rbp),%eax	
1210: c9	leaveq		
1211: c3	retq		

Dead Code: After (-02)

00000000000011b0 <dead_code>:

11b0: 8d 47 01

11b3: c3

lea 0x1(%rdi),%eax

retq

GCC Optimizations

- Constant Folding
- Common Sub-expression Elimination
- Dead Code
- **Strength Reduction**
- Code Motion
- Tail Recursion
- Loop Unrolling

Strength Reduction

Strength reduction changes divide to multiply, multiply to add/shift, and mod to AND to avoid using instructions that cost many cycles (multiply and divide).

```
int a = param2 * 32;
```

```
int b = a * 7;
```

```
int c = b / 2;
```

```
int d = param2 % 2;
```

```
for (int i = 0; i <= param2; i++) {
```

```
    c += param1[i] + 0x107 * i;
```

```
}
```

```
return c + d;
```

Shifting into Shifts

- `int a = param2 * 32;`

Becomes:

- `int a = param2 * 32;`

- `int b = a * 7;`

Becomes:

- `int b = a + (a << 2) + (a << 1);`

- `int c = b / 2;`

Becomes

- `int c = b >> 1`

GCC Optimizations

- Constant Folding
- Common Sub-expression Elimination
- Dead Code
- Strength Reduction
- **Code Motion**
- Tail Recursion
- Loop Unrolling

Code Motion

Code motion moves code outside of a loop if possible.

```
for (int i = 0; i < n; i++) {  
    sum += arr[i] + foo * (bar + 3);  
}
```

Common subexpression elimination deals with expressions that appear multiple times in the code. Here, the expression appears once, but is calculated each loop iteration, even though none of its values change during the loop.

Code Motion

Code motion moves code outside of a loop if possible.

```
int temp = foo * (bar + 3);  
for (int i = 0; i < n; i++) {  
    sum += arr[i] + temp;  
}
```

Moving it out of the loop allows the computation to happen only once.

Practice: GCC Optimization

```
int char_sum(char *s) {  
    int sum = 0;  
    for (size_t i = 0; i < strlen(s); i++) {  
        sum += s[i];  
    }  
    return sum;  
}
```

What is the bottleneck? What (if anything) can GCC do?

Practice: GCC Optimization

```
int char_sum(char *s) {  
    int sum = 0;  
    for (size_t i = 0; i < strlen(s); i++) {  
        sum += s[i];  
    }  
    return sum;  
}
```

What is the bottleneck? What (if anything) can GCC do?

strlen is called every loop iteration – code motion can pull it out of the loop

Tail Recursion

Tail recursion is an example of where GCC can identify recursive patterns that can be more efficiently implemented iteratively.

```
long factorial(int n) {  
    if (n <= 1) {  
        return 1;  
    }  
    else return n * factorial(n - 1);  
}
```

GCC Optimizations

- Constant Folding
- Common Sub-expression Elimination
- Dead Code
- Strength Reduction
- Code Motion
- Tail Recursion
- **Loop Unrolling**

Loop Unrolling

Loop Unrolling: Do **n** loop iterations' worth of work per actual loop iteration, so we save ourselves from doing the loop overhead (test and jump) every time, and instead incur overhead only every n-th time.

```
for (int i = 0; i <= n - 4; i += 4) {  
    sum += arr[i];  
    sum += arr[i + 1];  
    sum += arr[i + 2];  
    sum += arr[i + 3];  
} // after the loop handle any leftovers
```