

# Processes II & Virtual Memory I



# Lecture Topics

- **Processes and context switching**
  - **Creating new processes**
    - **fork() and exec\*()**
  - Ending a process
    - exit(), wait(), waitpid()
    - Zombies
- Virtual Memory (VM)
  - Overview and motivation
  - VM as a tool for caching
  - Address translation
  - VM as a tool for memory management
  - VM as a tool for memory protection


# fork Example

```
void fork1() {  
    int x = 1;  
    pid_t fork_ret = fork();  
    if (fork_ret == 0)  
        printf("Child has x = %d\n", ++x);  
    else  
        printf("Parent has x = %d\n", --x);  
    printf("Bye from process %d with x = %d\n", getpid(), x);  
}
```

- Both **parent** and **child** start/continue execution after fork
- **Child** gets a *copy* of **parent's** data - both processes start with `x = 1`
  - Subsequent changes to `x` are **independent**
- Shared open files - `stdout` is the same for both
- **Can't predict execution order** of **parent** and **child** - up to the OS!

# fork Example: Possible Output

```
void fork1() {  
    int x = 1;  
    pid_t fork_ret = fork();  
    if (fork_ret == 0)  
        printf("Child has x = %d\n", ++x);  
    else  
        printf("Parent has x = %d\n", --x);  
    printf("Bye from process %d with x = %d\n",  
        getpid(), x);  
}
```



# Polling Question

Which of the two sequences of outputs are possible?

```
void nestedfork() {  
    printf("L0\n");  
    if (fork() == 0) {  
        printf("L1\n");  
        if (fork() == 0) {  
            printf("L2\n");  
        }  
    }  
    printf("Bye\n");  
}
```

**Seq1**

L0

L1

Bye

Bye

Bye

L2

**A) No**

**B) No**

**C) Yes**

**D) Yes**

**Seq2**

L0

Bye

L1

L2

Bye

Bye

**No**

**Yes**

**No**

**Yes**

# Fork-Exec


- `fork()` creates a copy of the current process
- `exec*()` replaces the current process' code and address space with the code for a different program
  - Whole family of `exec` calls – see `exec(3)` and `execve(2)`

```
void fork_exec(char* path, char* argv[]) {  
    pid_t fork_ret = fork();  
    if (fork_ret != 0) {  
        printf("Parent: created a child %d\n", fork_ret);  
    } else {  
        printf("Child: about to exec a new program\n");  
        execv(path, argv);  
    }  
    printf("This line printed by parent only!\n");  
}
```

# exec-ing a Program

Very high-level diagram of what happens when you run the command “ls” in a Linux shell

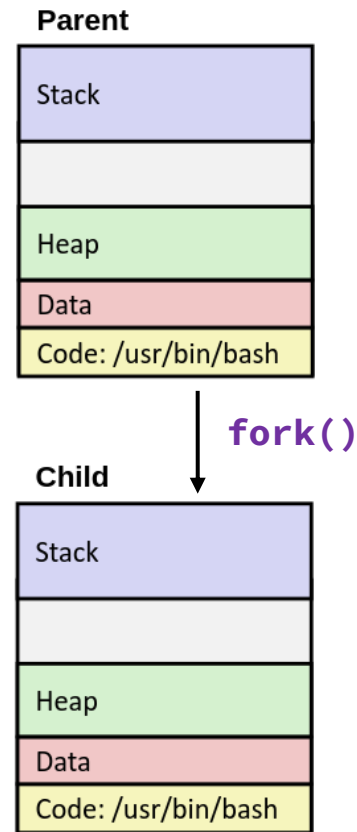
- This is the loading part of CALL!



## exec-ing a Program (pt 2)

Very high-level diagram of what happens when you run the command “ls” in a Linux shell

- This is the loading part of CALL!

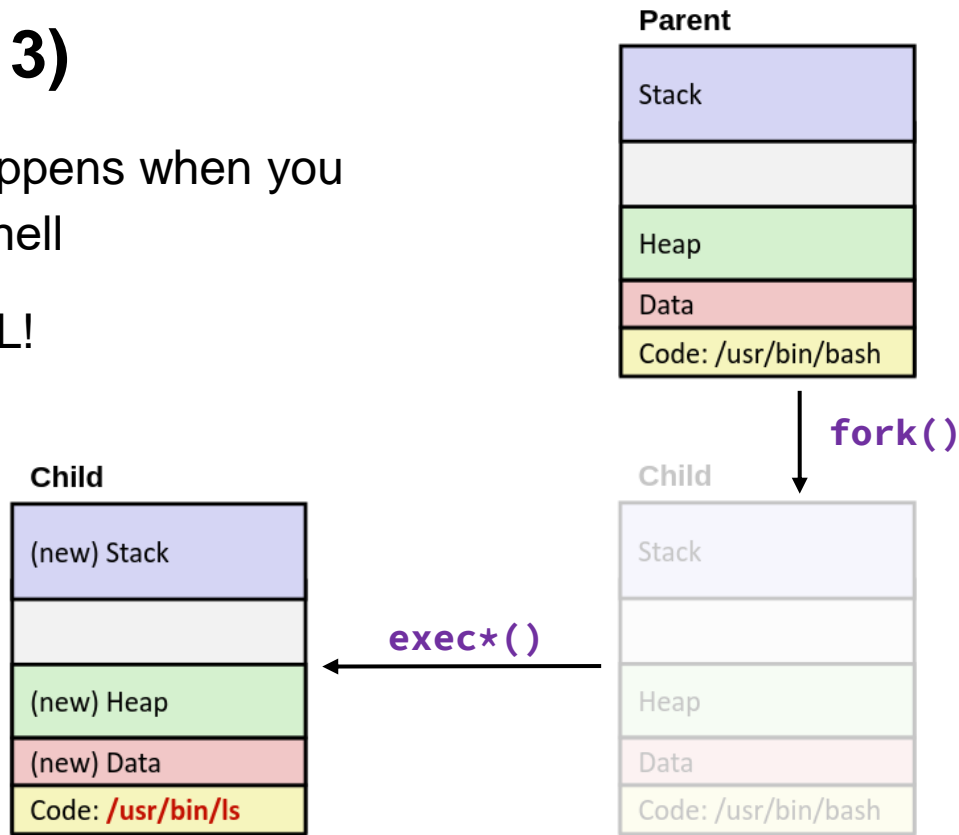




# exec-ing a Program (pt 3)

Very high-level diagram of what happens when you run the command “ls” in a Linux shell

- This is the loading part of CALL!



# Lecture Topics

- **Processes and context switching**
  - Creating new processes
    - `fork()` and `exec*()`
  - **Ending a process**
    - `exit()`, `wait()`, `waitpid()`
    - **Zombies**
- Virtual Memory (VM)
  - Overview and motivation
  - VM as a tool for caching
  - Address translation
  - VM as a tool for memory management
  - VM as a tool for memory protection

# exit: Exiting a Process

- `void exit(int status)`
  - Explicitly exits a process
    - Status code: 0 = normal exit, nonzero = abnormal exit
- The `return` statement from `main()` also exits a process
  - The return value is the status code
- Terminated processes still take up system resources
  - Data structures maintained by the OS
  - A process can't clean up all of its own resources when it exits, so whose responsibility is it?

# Zombies

- A terminated process that is still consuming resources is called a **zombie**
- **Parent** needs to **reap** its zombie **children** (i.e. clean up its resources)
  - **Parent** is given exit status information, then transfers control to the OS to delete zombie process
- What if the parent exits before reaping the child?
  - Orphaned child is reaped by `init` process (process 1)
    - Note: on recent Linux systems, `init` has been renamed to `systemd`




# wait: Synchronizing with Children

- `int wait(int* child_status)`
  - Suspends the current process until one of its children terminates
    - Reaps that child, then returns its PID
  - If `child_status != NULL`, then the `*child_status` value indicates why the child process terminated
    - If `NULL`, that means the status was ignored
    - Special macros for interpreting this status – see `man wait(2)`
- Note: If parent process has multiple children, `wait` will return when *any* of the children terminates
  - `waitpid` can be used to wait on a specific child process

# wait Example

```
void fork_wait() {  
    int child_status;  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```



Feasible output:

HC  
HP  
CT  
Bye

Infeasible output:

HP  
CT  
Bye  
HC

## wait Example 2: Zombies

```
void fork7() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Terminating Child, PID = %d\n",  
            getpid());  
        exit(0);  
    } else {  
        /* Parent */  
        printf("Running Parent, PID = %d\n",  
            getpid());  
        while (1); /* Infinite loop */  
    }  
}
```

Need to kill parent for  
init to reap the child

Zombie child is still  
there

```
linux> ./forks 7 &  
[1] 6639  
Running Parent, PID = 6639  
Terminating Child, PID = 6640  
linux> ps  
  PID TTY          TIME CMD  
 6585 tttyp9      00:00:00 tcsh  
 6639 tttyp9      00:00:03 forks  
 6640 tttyp9      00:00:00 forks <defunct>  
 6641 tttyp9      00:00:00 ps  
linux> kill 6639  
[1] Terminated  
linux> ps  
  PID TTY          TIME CMD  
 6585 tttyp9      00:00:00 tcsh  
 6642 tttyp9      00:00:00 ps
```

# wait Example 3: Non-terminating Child

```
void fork8() {  
    if (fork() == 0) {  
        /* Child */  
        printf("Running Child, PID = %d\n",  
            getpid());  
        while (1); /* Infinite loop */  
    } else {  
        /* Parent */  
        printf("Running Parent, PID = %d\n",  
            getpid());  
        exit(0);  
    }  
}
```

Child still active  
after parent  
terminates

```
linux> ./forks 8  
Terminating Parent, PID = 6675  
Running Child, PID = 6676  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9        00:00:00 tcsh  
 6676 ttyp9        00:00:06 forks  
 6677 ttyp9        00:00:00 ps  
linux> kill 6676  
linux> ps  
  PID TTY          TIME CMD  
 6585 ttyp9        00:00:00 tcsh  
 6678 ttyp9        00:00:00 ps
```

Must explicitly kill the  
child, or it will run forever.



# Lecture Topics


- Processes and context switching
  - Creating new processes
    - `fork()` and `exec*()`
  - Ending a process
    - `exit()`, `wait()`, `waitpid()`
    - Zombies
- Virtual Memory (VM\*)
  - Overview and motivation
  - VM as a tool for caching
  - Address translation
  - VM as a tool for memory management
  - VM as a tool for memory protection

**Warning:** Virtual memory is pretty complex, but crucial for understanding how processes work and for debugging performance

*\*Not to be confused with Virtual Machine, which is a whole other thing*



# Memory as we know it so far... is *virtual*!

- Programs refer to **virtual** memory addresses
  - System provides private addresses for each process
- Allocation: compiler and run-time system
  - Where different program objects should be stored
  - All allocation within single **virtual address space**
- But...
  - We *probably* don't have  $2^w$  bytes of physical memory
  - We *definitely* don't have  $2^w$  bytes of physical memory **for every process**
  - Processes should not interfere with each other
    - Except for specific cases where they want to share code or data



# Problem 1: How does everything fit?

64-bit virtual addresses can address 18 exabytes  
(18,446,744,073,709,551,616 bytes)



Physical main memory offers a few  
gigabytes  
(e.g., 8,589,934,592 bytes)



*(Not to scale; physical memory would be smaller  
than the period at the end of this sentence  
compared to the virtual address space.)*

# Problem 2: Memory Management

**We have multiple processes:**

Process 1  
Process 2  
Process 3  
...  
Process n

×

**Each process has:**


Stack  
Heap  
Static Data  
Literals  
Code  
...



Physical Memory




## Problem 3: How to protect data?



*What if two running programs both use the same address in their code?*

We want to make sure processes don't access the same physical memory locations.

## Problem 4: How to share data?




... Except sometimes we *do* want them to share memory!

- Inter-process communication
- Shared code
- etc.


# How can we solve these problems?

- “Any problem in computer science can be solved by adding another level of **indirection**.” – David Wheeler, inventor of the subroutine

Without indirection:




With indirection:



# Indirection

- The ability to reference something using a name, reference, or container instead of the value itself.
  - A flexible mapping between a name and a thing allows changing the thing without notifying holders of the name.
  - Adds some work (now have to look up 2 things instead of 1)
  - But don't have to track all uses of name/address (single source!)
- Examples:
  - Phone system: cell phone number portability
  - Domain Name Service (DNS): translation from name to IP address
  - Call centers: route calls to available operators, etc.

# Indirection in Virtual Memory




- Each process gets its own private address space
  - Translates to some location in physical memory
- Solves previous problems!



# Mapping

- A **virtual address** (VA) can be mapped to either **physical memory** (RAM) or on **disk**
  - Unused VAs may not have a mapping
  - VAs from *different* processes may (or may not) map to the same location in memory/disk



# Address Spaces


- **Virtual Address Space:** Set of  $N = 2^n$  virtual addresses
  - $\{0, 1, \dots, N-1\}$
  - Corresponds to word size (so in x86-64,  $n = 64$ )
- **Physical Address Space:** Set of  $M = 2^m$  physical addresses
  - $\{0, 1, \dots, M-1\}$
  - Address length  $m$  depends on hardware
- Every byte in main memory has:
  - **One** physical address (PA)
  - Zero, one, or *more* virtual addresses (VAs)

# Review Questions

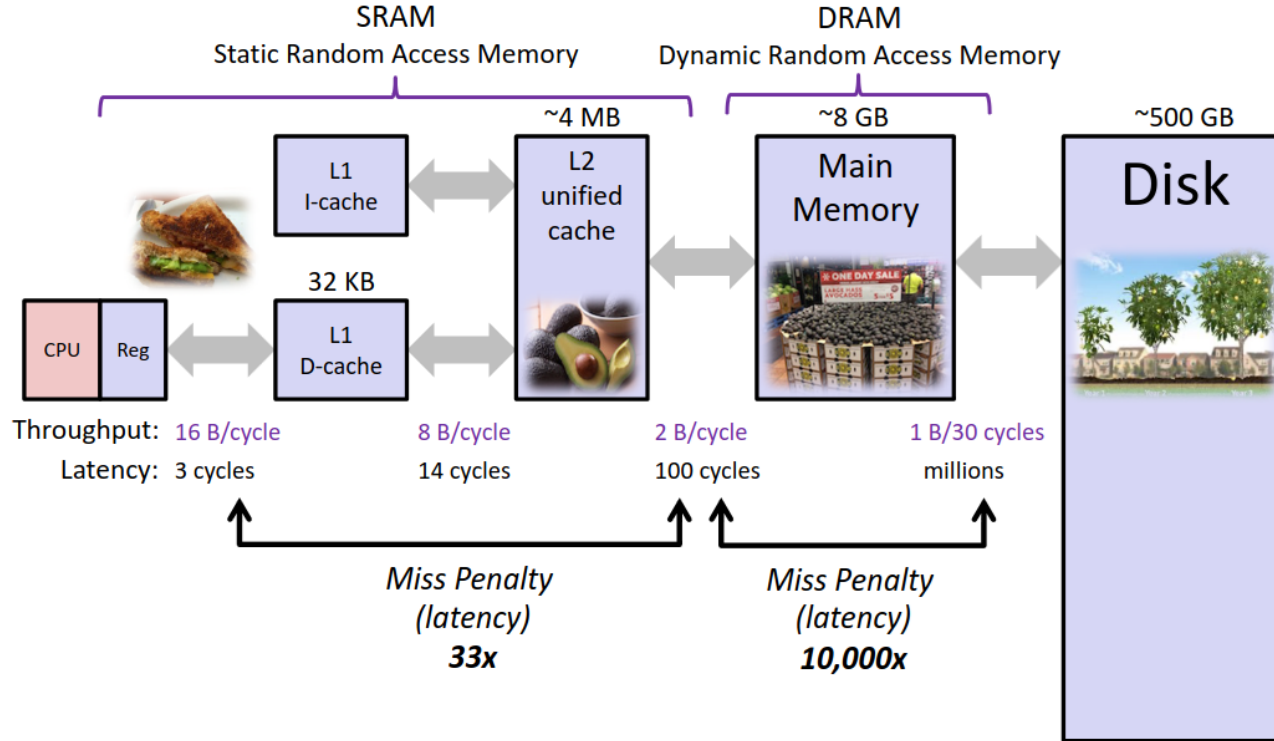
1. On a 64-bit machine currently running 8 processes, how much virtual memory is there?
2. True or False: A 32-bit machine with 8 GiB of RAM installed would never use all of it (in theory).

# VM and the Memory Hierarchy

- Think of memory (virtual or physical) as an array of bytes, now split into **pages**
  - Pages aligned (size is  $P = 2^p$  bytes), similar to cache blocks
  - Each virtual page can be stored in any physical page (no fragmentation!)
- Pages of virtual memory are usually stored in physical memory, but spill to disk when we run out of space
  - *Kind of like a cache!*



# Memory Hierarchy: Core 2 Duo



# Virtual Memory Design Consequences

- **Large page size:** typically 4-8 KiB or 2-4 MiB
  - Can be up to 1 GiB (for “Big Data” apps on big computers)
  - Much larger than cache blocks
- **Fully associative**
  - Any virtual page can be placed in any physical page
- Highly sophisticated, expensive **replacement algorithms** in OS
  - Too complicated and open-ended to be implemented in hardware
- **Write-back** rather than write-through
  - *Really* don't want to write to disk every time we modify memory
  - Some things may never end up on disk (e.g., stack for short-lived process)

# Why does VM work on RAM/disk?

- Avoids disk accesses because of *locality*
  - Same reason that L1 / L2 / L3 caches work
- The set of virtual pages that a program is “actively” accessing at any point in time is called its **working set**
  - If (working set of one process  $\leq$  physical memory):
    - Good performance for one process (after compulsory misses)
  - If (working sets of all processes  $>$  physical memory):
    - **Thrashing**: Performance meltdown where pages are swapped between memory and disk continuously
- *This is why your computer can feel faster when you add RAM*

# Summary

- **fork** makes two copies of the same process (**parent** & **child**)
  - Returns different values to the two processes
- **exec\*** replaces current process from file (new program)
- **exit** or return from `main` to end a process
- **wait** or **waitpid** used to synchronize **parent/child** execution and to **reap child**
- **Virtual memory** provides:
  - Ability to use limited memory (RAM) across multiple processes
  - Illusion of contiguous virtual address space for each process
  - Protection and sharing amongst processes